




Procedural Generation of Solvable Cooperative Levels for the Laser Learning Environment

Hugo Charels — Master thesis defense

Supervisors: Tom Lenaerts · Yannick Molinghen

ULB — Academic year 2025-2026 · 18/06/2026



Context & motivation

Laser Learning Environment (LLE)


- 2D grid-based fully cooperative multi-agent puzzle
- Each agent must reach its exit tile **simultaneously** with all others
- Laser beams are passable only by an agent of the **matching color**
- Reward signal is sparse: zero credit for intermediate coordination steps

The LLE cooperative mechanic


- A beam blocks **other** colours; each agent is **immune** to its own.
- Block your own beam → free a teammate's path (**locally unrewarded**).

Objective

Generate LLE levels **solvable** and that **require cooperation**.



LLE Level 6 – the hard target



Part 1 – Verification (RQ1–2)

How can we formally verify that a level is solvable and that solving it genuinely requires cooperation?

In this part we present our SAT-based verification tool:

RQ1 – Bounded-horizon solvability


- Formalise the decision problem
- Reduce to SAT: level + horizon $\rightarrow$ Boolean formula
- Verify solvability with a SAT solver

RQ2 – Cooperation detection

- Reuse the SAT encoding, change one rule
- Detect whether cooperation is required
- Classify cooperation structures

Key idea

SAT solver as a verification oracle: given a level and a horizon $T_{\max}$, answer “solvable?” in bounded time.



Bounded-horizon solvability $\rightarrow$ SAT

Decision problem

Given a level L and a horizon $T_{\max}$, does a valid joint trajectory exist?


SAT reduction

Build a Boolean formula $\Phi(L, T_{\max})$. Satisfiable $\Leftrightarrow$ solvable within $T_{\max}$.

$(L, T_{\max}) \rightarrow$ encode Φ $\rightarrow$ SAT solver $\rightarrow$ SAT / UNSAT $\rightarrow$ solvable / unsolvable

Why SAT? Fast, well-known problem, modern solvers are highly efficient.

Complexity: LLE solvability $\in$ NP and $\leq_p$ SAT. NP-hardness **open**.



Detecting cooperation


Idea: keep the SAT encoding, change **one** rule: beams no longer block same-colour agents → the cooperation mechanic is **disabled**.

Cooperation criterion


L requires cooperation with $T_{\max}$ $\iff \Phi(L, T_{\max})$ **SAT** and $\Phi_{\text{strict}}(L, T_{\max})$ **UNSAT**.

Three possible outcomes:

Condition	Label	Meaning
Φ UNSAT	UNSOLVABLE	No valid trajectory exists
Φ_{strict} UNSAT	COOPERATIVE	Every solution blocks a beam – cooperation required
Otherwise	NON_COOPERATIVE	A solution exists without the blocking mechanic



Horizon matters




Horizon	Solvable	Coop.
$T_{\max} \leq 2$	no	n/a
$3 \leq T_{\max} \leq 8$	yes	yes
$T_{\max} \geq 9$	yes	no

- **Too small** → agents run out of moves; the level is falsely marked unsolvable
- **Too large** → agents can take longer detours that bypass cooperation requirements

Trade-off


Larger $T_{\max}$ → more actions (more solutions), but bigger formula (longer solve duration). Pick the **smallest** $T_{\max}$ that still admits a solution.



Cooperation profiles

From one SAT model $\rightarrow$ a **dependency graph (edges: helper $\rightarrow$ beneficiary)**

classify by priority: fully coupled $\succ$ mutual $\succ$ distributed $\succ$ chain $\succ$ asymmetric.





Cooperation profiles

From one SAT model → a **dependency graph** (edges: helper → beneficiary)

classify by priority: fully coupled $\succ$ mutual $\succ$ distributed $\succ$ chain $\succ$ asymmetric.

asymmetric




Cooperation profiles



From one SAT model → a **dependency graph** (edges: helper → beneficiary)

classify by priority: fully coupled $\succ$ mutual $\succ$ distributed $\succ$ chain $\succ$ asymmetric.

asymmetric



chain




Cooperation profiles


From one SAT model → a **dependency graph** (edges: helper → beneficiary)

classify by priority: fully coupled $\succ$ mutual $\succ$ distributed $\succ$ chain $\succ$ asymmetric.





asymmetric



chain



distributed




Cooperation profiles



From one SAT model → a **dependency graph** (edges: helper → beneficiary)

classify by priority: fully coupled $\succ$ mutual $\succ$ distributed $\succ$ chain $\succ$ asymmetric.



asymmetric





chain



distributed



mutual




Cooperation profiles


From one SAT model → a **dependency graph** (edges: helper → beneficiary)

classify by priority: fully coupled $\succ$ mutual $\succ$ distributed $\succ$ chain $\succ$ asymmetric.


asymmetric




chain










distributed



mutual



fully coupled



Part 2 – Construction (RQ3–4)

How can we embed formal verification inside a procedural generator so that every accepted level is certified, and can we control the cooperation structure by targeting specific profiles?

In this part we present our level generators:

RQ3 – Generator efficiency


- Random placement with geometric constraints
- Constructive lane-based approach
- Level-6-style clustered design
- SAT oracle filters solvable / cooperative levels

RQ4 – Output diversity


- Five-profile taxonomy of cooperation structures
- Generators can steer towards specific profiles
- Constructive → mutual by design

Key idea

Place-and-check with a SAT oracle: every accepted level is **provably** solvable.




Random Generator




Pure random

Lasers point off-grid, a wall block a laser source




Random Generator



Pure random

Lasers point off-grid, a wall block a laser source




Constrained random

Filters out laser sources that behave like walls.




Constructive Cooperative Generator

Oriented **lane-based** design. Each agent gets a reserved lane, cooperation arises from crossing lasers. Near-zero rejection rate.




1. Reserve lanes, place agents + exits




Constructive Cooperative Generator


Oriented **lane-based** design. Each agent gets a reserved lane, cooperation arises from crossing lasers. Near-zero rejection rate.



1. Reserve lanes, place agents + exits




2. Place walls in free cells




Constructive Cooperative Generator


Oriented **lane-based** design. Each agent gets a reserved lane, cooperation arises from crossing lasers. Near-zero rejection rate.




1. Reserve lanes, place agents + exits



2. Place walls in free cells




3. Add lasers crossing lanes




Level-6-Style Generator

Clustered starts + exits on opposite sides. Shared laser corridor = mutual by design. SAT verifies.




1. Place clusters on opposite sides




Level-6-Style Generator

Clustered starts + exits on opposite sides. Shared laser corridor = mutual by design. SAT verifies.



1. Place clusters on opposite sides




2. Fill the central corridor with lasers




Level-6-Style Generator


Clustered starts + exits on opposite sides. Shared laser corridor = mutual by design. SAT verifies.




1. Place clusters on opposite sides



2. Fill the central corridor with lasers




3. Add walls to shape the solution space



Generator efficiency (RQ3)


Rejection rate by generator and grid size



- 3x3 / 2-agents / 1-laser
- 5x5 / 3-agents / 2-lasers
- 8x8 / 4-agents / 3-lasers



Output diversity (RQ4)



- Constructive → 92% mutual.
- Level-6-style → 68% mutual.
- Random → mostly asymmetric.

(8x8 / 3-agents / 2-lasers)



Part 3 — Experiment (RQ5–6)

Can MARL agents trained on certified levels generalise to held-out ones, and does a staged curriculum outperform direct training on a hard target?

In this part we investigate whether generated levels are learnable:

RQ5 — Learnability


- Train IQL / VDN / QMIX on certified levels
- Measure train vs. test success
- Scale training data to close the generalisation gap

RQ6 — Curriculum learning

- Compare curriculum orderings (direct, forward, mixed, reverse)
- Test harder cooperation targets (mutual, Level-6)




Learnability (RQ5)



(200k steps, 5x5 / 2-agents / 1-laser, asymmetric)

Takeaway

Learnable at 5×5, but too few levels $\Rightarrow$ overfitting



Scaling data (RQ5)



(300k steps, 5x5 / 2-agents / 1-laser, asymmetric)


Takeaway

Unlimited certified levels $\Rightarrow$
overfitting becomes
generalisation.




Curriculum ordering (RQ6)

Curriculum strategy on the 6x6/1L target (n=8 seeds, 95% CI)



Four budget-matched schedules (400k steps, 6×6 / 2-agents / 1-laser, asymmetric)



Curriculum targets (RQ6)

Target	Grid	Coop.	Budget	Test
Learnability	5×5	asym	200k	0.20
Curric. order	6×6	asym	400k	0.17
Frontier	6×6	mutual	600k	0.00
Frontier	5×5	mutual	200k	0.00
Level-6	12×13	mutual	2M	0.00

- Reachable target: ordering doesn't beat **direct**; only data **diversity** helps.
- A **cliff** at **mutual** cooperation: success → **zero**, even at 10× budget.

Takeaway

Scoped **negative** result: the base **mutual** task is unlearnable by IQL/VDN/QMIX.



Beyond the thesis: upstream

Solver, cooperation detector, and generators are **merged into the official LLE library** (`laser-learning-environment[generator]`, PyPI, v2.9.0+).

Public API

```
lle.solve • lle.is_cooperative • lle.cooperation_level • lle.generate
```

Takeaway

Not just thesis artefacts — **usable tooling** for the LLE users.



Conclusion

- RQ1 ✓ – SAT verification tool for solvability; in NP, $\leq_p$ SAT.
- RQ2 ✓ – strict counterfactual: $\text{coop.} \Leftrightarrow \text{SAT} \wedge \text{strict-UNSAT}$.
- RQ3 ✓ – SAT-as-oracle generators; every accepted level **certified**.
- RQ4 ✓ – five-profile taxonomy; generators steer toward **mutual**.
- RQ5 ✓ – certified levels learnable; scaling data closes the gap.
- RQ6 ✗ – value-based MARL can't learn **mutual** cooperation; no curriculum helps.

Takeaway

Central contribution: procedural generation coupled to **explicit certification**.



Future work

Theory & model

- Settle **NP-hardness** of LLE solvability.
- Add **gems** and true **void** tiles.
- Richer (time-indexed / weighted) cooperation graphs.

The certification framework is settled; the open question is **how far it extends**.

Learning

- **Make mutual coordination learnable:** dense / intrinsic rewards.
- Deeper research into cooperation structures: how do they affect learnability? Is the learning of one profile a stepping to another?
- Beyond value-based: actor-critics (MADDPG, COMA), MAVEN.




Summary

Contributions

1. SAT verification tool for LLE solvability.
2. Cooperation detector + five-profile taxonomy.
3. Solver-in-the-loop certified generators.
4. Merged into the official LLE library.

Key results

1. Constructive generator $\approx$ free.
2. Scaling certified data closes the generalisation gap.
3. A learnability **cliff** at mutual cooperation.



Backup slides